

+918668416372

INDIA

Ksohail0321@gmail.com

KHAN SOHIL FEROZ

DATA ENGINEER | PYTHON DEVELOPER | KAFKA DEVELOPER & ADMIN |

Objective

Results-driven Data Engineer with over 4 years of experience specializing in Big Data and Kafka ecosystem management Pyspark and Databricks. Proven expertise in optimizing and implementing robust security measures, and deploying production-grade systems. Proficient in Python for data ingestion, extraction, and Kafka producer/consumer development. Skilled in DevOps practices, including GIT , Docker and kubernetes.

Professional Experience:

DURATION: Dec 2020 – April 2024

Designation: Data Engineer.

Previous company: DREAMCARE DEVELOPERS PVT. LTD.

DURATION: April 2024 – December 2024.

Designation: Solution Developer.

Current company: Tata Technologies.

DURATION: February 2025 - Till Date

Designation: Software Engineer.

Current company: Saint-Gobain.

Roles & Responsibilities:

- Driven and passionate Kafka Administration and Development professional with over 4 years of hands-on experience in installing, configuring, supporting, and overseeing Big Data Clusters. Proficient in managing various distributions such as Cloudera Enterprise.
- Proficient in data ingestion and extraction using Python and libraries such as pandas, boto3, pycopg2.
- Implemented custom Serializers and Deserializers for efficient message serialization using JSON.
- Built and maintained pipelines for streaming data from heterogeneous sources to cloud storage and downstream systems in python.
- Integrated Kafka with cloud storage services (S3) for seamless data ingestion and archival.

- Spearheading enterprise-scale Kafka modernization by integrating **Kafka Connect**, Debezium, and Confluent Cloud to build secure, scalable event pipelines.
- Built streaming ETL pipelines using ksqlDB and PySpark **on** Databricks, integrating data from PostgreSQL, MongoDB Atlas, and Kafka topics.
- Architected ingestion pipelines from Kafka to Delta Lake and Snowflake, using auto-ingest and batch processing for low-latency analytics.
- Designed and implemented stateful stream processing prototypes using Apache Flink, working on sessionization and fraud detection POCs.
- Deployed and managed Kafka Connect clusters with custom JARs and SMTs for HTTP Sink and JDBC connectors.
- Collaborated closely with product and BI teams to ensure event schema consistency, performance tuning, and real-time business KPIs.
- Actively contributed to internal capability building on ksqlDB, Delta Lake, and Kafka security (SASL/SSL).
- Designed and implemented Kafka producers and consumers using Python, ensuring efficient data flow within the Kafka ecosystem.
- Improved system reliability and efficiency by introducing automation python scripts for routine Kafka maintenance tasks.
- Implemented CDC (Change Data Capture) for real-time data synchronization between relational databases and Kafka topics, ensuring minimal latency and consistency across systems.
- Monitored and maintained CDC connectors for real-time change capture, using Kafka metrics and monitoring tools to ensure smooth operations
- Worked on Streaming data using Nifi and Kafka from different sources.
- Experienced in handling AWS services like IAM, EC2, S3, VPC, Cloudwatch.
- Configured and managed the Backup and Disaster Recovery for Big Data in Production.
- Troubleshooting and checking Log files and providing optimal solutions in timely manner.
- Experienced in planning the Capacity of Big Data Cluster according to the requirements.
- Experienced in Data Integration and Data Migration.
- In depth knowledge of AWS, including cluster setup, configuration and optimization.
- Experienced in Implementing Data Lake and Data Warehouse using Cloudera Enterprise.
- Perform Root Cause Analysis (RCA) for failures and provide optimal solution within SLA's.
- Monitor application workflows and perform periodic trend analysis to proactively identify performance optimization opportunities.
- Develop and Deliver Reports of Big Data Kafka Cluster.
- Decision making/problem solving skills.
- Coordinate with fellow teams including development, QA, infrastructure, network and database to achieve effective resolution to complex problems.
- Followed standard Back up policies.
- Experience In Performance Tuning
- Followed Best Practices for Kafka Cluster.
- Exposure Of Devops Tools Like Git, Jenkins, Docker.
- Experienced in Implementing Docker.
- Experience in Implementing CI/CD Pipeline From GitHub to Jenkins.
- Experience in Implementing Apache Kafka.
- Experience in Implementing Cruise Control and Streams Messaging Manager.

- Experience in Monitoring and Troubleshooting of Kafka Using Streams Messaging Manager and Cruise Control.
- Administered and maintained Kafka clusters, ensuring high availability and reliability.
- Managed Kafka topics, partitions, and brokers for optimal performance and scalability.
- Implemented and monitored Kafka security features.
- Worked closely with development teams to troubleshoot and optimize Kafka producers and Consumers.
- Implemented Kafka best practices for data retention, compaction, and log segmentation.
- Experience in Trouble Shooting Kafka.
- Experience in Implementing Apache and Confluent Kafka.
- Designed and implemented Kafka producers and consumers using Python and Connectors for seamless data flow.
- Configured Kafka Connect with source and sink connectors.
- Optimized Kafka clusters, topics, partitions, and brokers for high availability and performance.
- Configured and maintained CDC pipelines using Debezium and Confluent Kafka for real-time data replication.
- Implemented SSL/TLS encryption and Kerberos authentication for securing Kafka clusters and ensuring compliance with organizational security standards.
Optimized consumer fetch configurations, including `fetch.min.bytes` and `max.poll.records`, for efficient data consumption.
- Improved partitioning strategies by implementing custom partitioners for better data distribution and load balancing.
- Tuned Kafka producer configurations, including batch size, `linger.ms`, and compression settings, to enhance throughput and reduce latency.
- Integrated Kafka with external systems using Kafka Connect.
- Built real-time data enrichment pipelines by consuming data from multiple Kafka topics and enriching it with external API data.
- Ensuring data integrity and system reliability.
- Experience in Implementing Apache Nifi.
- Implemented Kubernetes (K8s) for production-grade systems with a 3-node master-worker architecture.

Worked on Security:

- Worked on Big Data security :
- For Authentication used Kerberos – MIT & AD
- For Authorization used Extended ACLs, Posix Style Permission, Sticky Bits, Ranger.
- Encryption for Data in Motion used SSL, TLS & for Data at Rest used Encryption Zone.

Skills:

Big Data Technologies : Apache Kafka ,Confluent Kafka,Apache Spark
DevOps : Git, Jenkins, Docker.
Eco System : HDFS, Yarn, Hive, Hue, Sqoop, Zookeeper,Kafka, Spark .
Cloud Platform : AWS, EC2, S3,Cloudwatch, VPC, IAM.
Security : MIT Kerberos,KMS.
Containerization : Kubernetes
Programming : Python,SQL.
AI Tools : ChatGPT, Gemini, GitHub Copilot.

Education:

B.E(Computer Science and Engineering): Marathwada Institute of Technology, Aurangabad.

Personal Strength: oneself, intrinsic motivation, rapid learning ability, adaptability, and flexibility.